

# The basics of programming in C



# The basics of programming in C

```
for (;;) {  
    n = readlink(pathname, buff->mem, buff->cap - 1);  
    if (n != (ssize_t)(buff->cap - 1))  
        break;  
    string_buff_grow(buff);  
}
```

## The basics of programming in C

```
for (;;) {  
    n = readlink(pathname, buff->mem, buff->cap - 1);  
    if (n != (ssize_t)(buff->cap - 1))  
        break;  
    string_buff_grow(buff);  
}
```

Do not make things more complicated than needed.

Read `/src/linux/fs/namei.c:getname_flags()`. In Linux there is no way to create a symlink that is longer than 4096 (= `PATH_MAX`) bytes.

```
assert(close(currdir) >= 0);
```

## The basics of programming in C

```
assert(close(currdir) >= 0);
```

```
close(curdir);
```

Second, it is simpler.

```
assert(close(currdir) >= 0);
```

```
close(curdir);
```

Second, it is simpler.

First, `assert ( )` is a macros that can expand to an empty string.

```
% cat assert.c
```

```
#include <stdio.h>
```

```
#include <assert.h>
```

```
int main()
```

```
{
```

```
    assert(printf("hello, world\n") >= 0);
```

```
    return 0;
```

```
}
```

```
% gcc -DNDEBUG assert.c
```

```
% ./a.out
```

```
%
```

## The basics of programming in C

```
char* token = strtok(temp_path, "/");  
while (token != NULL) {  
    ...  
}
```

## The basics of programming in C

```
char* token = strtok(temp_path, "/");  
while (token != NULL) {  
    ...  
}
```

This is not thread-safe.

Glibc has `strtok_r()`, but it is better to avoid this family of functions altogether.



## The basics of programming in C

```
for (i = 0; i < QUEUE_DEPTH; i++) {  
    sqe = io_uring_get_sqe(&ring);  
    if (!sqe) {  
        ret = -EAGAIN;  
        goto cleanup;  
    }  
    ...  
}
```

## The basics of programming in C

```
for (i = 0; i < QUEUE_DEPTH; i++) {  
    sqe = io_uring_get_sqe(&ring);  
    if (!sqe) {  
        ret = -EAGAIN;  
        goto cleanup;  
    }  
    ...  
}
```

```
for (i = 0; i < QUEUE_DEPTH; i++) {  
    sqe = io_uring_get_sqe(&ring);  
    ... sqe is guaranteed to be non-NULL ...  
}
```

Don't make things more complex than needed.

`io_uring_get_sqe()` fails only if there are no more entries in the submission queue.

## The basics of programming in C

```
int blocks_count = file_sz / BLOCK_SZ;  
if (file_sz % BLOCK_SZ)  
    blocks_count++;
```

## The basics of programming in C

```
int blocks_count = file_sz / BLOCK_SZ;  
if (file_sz % BLOCK_SZ)  
    blocks_count++;
```

```
int blocks_count = (file_sz + BLOCK_SZ - 1) / BLOCK_SZ;
```

# The basics of programming in C

```
struct io_data {  
    char  *buf;  
    size_t size;  
    off_t offset;  
    int  read_done;  
};
```

...

```
for (i = 0; i < QUEUE_DEPTH; i++) {  
    data[i] = malloc(sizeof(struct io_data));  
    data[i]->buf = malloc(COPY_BLOCK_SIZE);  
    data[i]->offset = i * COPY_BLOCK_SIZE;  
    data[i]->read_done = 0;  
}
```

# The basics of programming in C

```
struct io_data {  
    char *buf;  
    size_t size;  
    off_t offset;  
    int read_done;  
};
```

...

```
for (i = 0; i < QUEUE_DEPTH; i++) {  
    data[i] = malloc(sizeof(struct io_data));  
    data[i]->buf = malloc(COPY_BLOCK_SIZE);  
    data[i]->offset = i * COPY_BLOCK_SIZE;  
    data[i]->read_done = 0;  
}
```

```
struct io_data {  
    off_t offset;  
    int size;  
    bool read_done;  
    char buf[];  
};
```

...

```
for (i = 0; i < QUEUE_DEPTH; i++) {  
    data[i] = xmalloc(sizeof(*data[i]) + COPY_BLOCK_SIZE);  
    data[i]->offset = i * COPY_BLOCK_SIZE;  
    data[i]->read_done = 0;  
}
```

This is called **flexible arrays**.

# The basics of programming in C

```
struct io_data {  
    char *buf;  
    size_t size;  
    off_t offset;  
    int read_done;  
};
```

...

```
for (i = 0; i < QUEUE_DEPTH; i++) {  
    data[i] = malloc(sizeof(struct io_data));  
    data[i]->buf = malloc(COPY_BLOCK_SIZE);  
    data[i]->offset = i * COPY_BLOCK_SIZE;  
    data[i]->read_done = 0;  
}
```

```
struct io_data {  
    off_t offset;  
    int size;  
    bool read_done;  
    char buf[] __attribute__((__element_count__(size)));  
};
```

...

```
for (i = 0; i < QUEUE_DEPTH; i++) {  
    data[i] = xmalloc(sizeof(*data[i]) + COPY_BLOCK_SIZE);  
    data[i]->offset = i * COPY_BLOCK_SIZE;  
    data[i]->read_done = 0;  
}
```

GCC 15 is going to have bound checks on flexible arrays.

**See also:** <https://lwn.net/Articles/908817/>

## The basics of programming in C

```
static char is_directory(const char *path) {  
    struct stat path_stat;  
  
    if (stat(path, &path_stat) == -1)  
        return 0;  
  
    return S_ISDIR(path_stat.st_mode);  
}
```

```
static char is_link(const char* child) {  
    ...  
}
```

...

```
char is_link_result = is_link(child);
```

```
if (is_link_result > 0) {  
    ...  
} else {  
    ...  
    if (is_directory(state.path)) {  
        append_dir();  
    }  
    ...  
}
```



## The basics of programming in C

```
static char is_directory(const char *path) {  
    struct stat path_stat;  
  
    if (stat(path, &path_stat) == -1)  
        return 0;  
  
    return S_ISDIR(path_stat.st_mode);  
}
```

```
static char is_link(const char* child) {  
    ...  
}
```

...

```
char is_link_result = is_link(child);
```

```
if (is_link_result > 0) {  
    ...  
} else {  
    ...  
    if (is_directory(state.path)) {  
        append_dir();  
    }  
    ...  
}
```

These system calls are redundant. Avoid this because a system call is an expensive operation. It used to be a context switch. Nowadays, it switches the context, switches the page tables, flushes caches, etc.

It suffices to `stat()` the current file only once.

# The basics of programming in C

```
struct dirwalk
```

```
{  
    char **components;  
    int len;  
    int cap;  
    int fd;  
};
```

```
int dirwalk_enter_one(struct dirwalk *d, const char *name)  
{  
    int fd = openat(d->fd, name, O_RDONLY|O_NOFOLLOW);  
    ... this is enough to detect symlinks and –ENOTDIR ...  
}
```

```
int dirwalk_enter_path(struct dirwalk_d, const char *path)  
{  
    ... split path and dirwalk_enter_one() ...  
}
```